

error[E0382]: borrow of moved value

Secure Systems Engineering Spring 2024

🛡️ EE G7701

February 13, 2024

Tushar Jois



Recap

- Buffer overflows can be used to hijack the control flow of a program
- Either shellcode or a return-to-libc can spawn a shell afterward
- Defenses against buffer overflows exist, but have shortcomings

Lesson objectives

- Become familiar with basic Rust syntax and types
- Apply Rust languages features when writing code
- Intuitively understand Rust's borrowing and ownership semantics

Assignment 1 reflection

- How do/did we feel about it?
 - It's OK if you don't have as much experience with code
 - But, it's still a graduate-level course
- Got a lot of emails from folks asking about how to do things
 - Emailing the professor is typically a last resort
- Problem solving is part of the assignment
 - Be specific about what you did and why you did it
 - Always start by closely doing all of the reading
 - Use online resources to help you learn
 - Help each other out
 - Work with your partner -- pair programming
 - Discussion board -- anonymously ask/give help
- Deadline extension -- Assignment 1 now due 10p Feb 14
 - One-time -- people were having some initial trouble with the VM
 - Don't ask for deadline extensions 6 hours before the due date
 - Not fair to those of your classmates who prioritized effectively

Attribution

This lecture is adapted from material developed for a different course, [CIS 198: Rust Programming](#), Spring 2016, at the University of Pennsylvania.

Rust?

“A language empowering everyone to build reliable and efficient software.”

- rust-lang.org

Rust is:

- Fast
- Safe
- Functional
- Zero-cost

Rust is.. fast

- Rust compiles to native code
- Rust has no garbage collector
- Most abstractions have zero cost
- Fine-grained control over lots of things
- Pay for exactly what you need...
- ...and pay for most of it at compile time

Rust is... safe

- No null
- No uninitialized memory
- No dangling pointers
- No double free errors
- No manual memory management!

Rust is... functional

- First-class functions
- Trait-based generics
- Algebraic datatypes
- Pattern matching

Rust is... zero-cost

- Zero-cost, 100% safe abstractions form Rust's defining feature
- Strict compile-time checks remove need for runtime
- Big concept: Ownership

```
fn main() {  
    println!("Hello, world!");  
}
```

Basic Rust Syntax

Variable Bindings

- Variables are bound with `let`:

```
let x = 17;
```

- Bindings are implicitly-typed: the compiler infers based on context.
- The compiler can't always determine the type of a variable, so sometimes you have to add type annotations.

```
let x: i16 = 17;
```

- Variables are inherently immutable:

```
let x = 5;
```

```
x += 1; // error: re-assignment of immutable variable x
```

```
let mut y = 5;
```

```
y += 1; // OK!
```

Variable Bindings

- Bindings may be shadowed:

```
let x = 17;  
let y = 53;  
let x = "Shadowed!";  
// x is not mutable, but we're able to re-bind it
```

- The shadowed binding for `x` above lasts until it goes out of scope.
- Above, we've effectively lost the first binding, since both `xs` are in the same scope.
- Patterns may also be used to declare variables:

```
let (a, b) = ("foo", 12);
```

Expressions

- (Almost!) everything is an expression: something which returns a value.
 - Exception: variable bindings are not expressions.
- The “nothing” type is called “unit”, which is written `()`.
 - The *type* `()` has only one value: `()`.
 - `()` is the default return type.
- Discard an expression’s value by appending a semicolon. Now it returns `()`.
 - Hence, if a function ends in a semicolon, it returns `()`.

```
fn foo() -> i32 { 5 }  
fn bar() -> () { () }  
fn baz() -> () { 5; }  
fn qux()      { 5; }
```

Expressions

- Because everything is an expression, we can bind many things to variable names:

```
let x = -5;  
let y = if x > 0 { "greater" } else { "less" };  
println!("x = {} is {} than zero", x, y);
```

- Aside: "{}" is Rust's (most basic) string interpolation operator
 - Similar to Python, Ruby, C#, and others; like printf's "%s" in C/C++.

Comments

```
/// Triple-slash comments are docstring comments.
///
/// `rustdoc` uses docstring comments to generate
/// documentation, and supports Markdown formatting.
fn foo() {
    // Double-slash comments are normal.

    /* Block comments
       * also exist /* and can be nested! */
       */
}
```


Types

Primitive Types

- `bool`: spelled `true` and `false`.
- `char`: spelled like `'c'` or `'🐱'` (chars are Unicode!).
- Numerics: specify the signedness and size.
 - `i8, i16, i32, i64, isize`
 - `u8, u16, u32, u64, usize`
 - `f32, f64`
 - `isize` & `usize` are the size of pointers (and therefore have machine-dependent size)
 - Literals are spelled like `10i8, 10u16, 10.0f32, 10usize`.
 - Type inference for non-specific literals default to `i32` or `f64`:
 - e.g. `10` defaults to `i32`, `10.0` defaults to `f64`.
- Arrays, slices, `str`, tuples.
- Functions.

Arrays

- Arrays are generically of type `[T; N]`.
 - `N` is a compile-time *constant*. Arrays cannot be resized.
 - Array access is bounds-checked at runtime.
 - Arrays are indexed with `[]` like most other languages:
 - `arr[3]` gives you the 4th element of `arr`
- ```
let arr1 = [1, 2, 3]; // (array of 3 elements)
let arr2 = [2; 32]; // (array of 32 `2`s)
```

# Slices

- Generically of type `&[T]`
- A “view” into an array by reference
- Not created directly, but are borrowed from other variables
- Mutable or immutable
- How do you know when a slice is still valid? Coming soon...

```
let arr = [0, 1, 2, 3, 4, 5];
let total_slice = &arr; // Slice all of `arr`
let total_slice = &arr[..]; // Same, but more explicit
let partial_slice = &arr[2..5]; // [2, 3, 4]
```

# Strings

- Two types of Rust strings: `String` and `&str`.
- `String` is a heap-allocated, growable vector of characters.
- `&str` is a type<sup>1</sup> that's used to slice into `Strings`.
- String literals like `"foo"` are of type `&str`.

```
let s: &str = "galaxy";
let s2: String = "galaxy".to_string();
let s3: String = String::from("galaxy");
let s4: &str = &s3;
```

<sup>1</sup>`str` is an unsized type, which doesn't have a compile-time known size, and therefore cannot exist by itself.

# Tuples

- Fixed-size, ordered, heterogeneous lists
- Index into tuples with `foo.0`, `foo.1`, etc.
- Can be destructured in `let` bindings

```
let foo: (i32, char, f64) = (72, 'H', 5.1);
```

```
let (x, y, z) = (72, 'H', 5.1);
```

```
let (a, b, c) = foo; // a = 72, b = 'H', c = 5.1
```

# Casting

- Cast between types with `as`:

```
let x: i32 = 100;
```

```
let y: u32 = x as u32;
```

- Naturally, you can only cast between types that are safe to cast between.
  - No casting `[i16; 4]` to `char`! (This is called a “non-scalar” cast)
  - There are unsafe mechanisms to overcome this, if you know what you’re doing.

# Vec<T>

- A standard library type: you don't need to import anything.
- A `Vec` (read “vector”) is a heap-allocated growable array.
  - (cf. Java's `ArrayList`, C++'s `std::vector`, etc.)
- `<T>` denotes a generic type.
  - The type of a `Vec` of `i32`s is `Vec<i32>`.
- Create `Vecs` with `Vec::new()` or the `vec!` macro.
  - `Vec::new()` is an example of namespacing. `new` is a function defined for the `Vec` struct.



# Vec<T>

*// Explicit typing*

```
let v0: Vec<i32> = Vec::new();
```

*// v1 and v2 are equal*

```
let mut v1 = Vec::new();
```

```
v1.push(1);
```

```
v1.push(2);
```

```
v1.push(3);
```

```
let v2 = vec![1, 2, 3];
```

*// v3 and v4 are equal*

```
let v3 = vec![0; 4];
```

```
let v4 = vec![0, 0, 0, 0];
```

# Vec<T>

```
let v2 = vec![1, 2, 3];
let x = v2[2]; // 3
```

- Like arrays, vectors can be indexed with [].
  - You can't index a vector with an i32/i64/etc.
  - You must use a `usize` because `usize` is guaranteed to be the same size as a pointer.
  - Other integers can be cast to `usize`:

```
let i: i8 = 2;
let y = v2[i as usize];
```
- Vectors have an extensive `stdlib` method list, which can be found at the [official Rust documentation](#).

# References

- Reference *types* are written with an `&`: `&i32`.
- References can be taken with `&` (like C/C++).
- References can be *dereferenced* with `*` (like C/C++).
- References are guaranteed to be valid.
  - Validity is enforced through compile-time checks!
- These are *not* the same as pointers!
- Reference lifetimes are pretty complex, as we'll explore later.

```
let x = 12;
let ref_x = &x;
println!("{}", *ref_x); // 12
```

# Control Flow

# If Statements

```
if x > 0 {
 10
} else if x == 0 {
 0
} else {
 println!("Not greater than zero!");
 -10
}
```

- No parens necessary.
- Entire if statement evaluates to one expression, so every arm must end with an expression of the same type.
  - That type can be unit ():

```
if x <= 0 {
 println!("Too small!");
}
```

# Loops

- Loops come in three flavors: `while`, `loop`, and `for`.
  - `break` and `continue` exist just like in most languages
- `while` works just like you'd expect:

```
let mut x = 0;
while x < 100 {
 x += 1;
 println!("{}", x);
}
```

# Loops

- `loop` is equivalent to `while true`, a common pattern.
  - Plus, the compiler can make optimizations knowing that it's infinite.

```
let mut x = 0;
loop {
 x += 1;
 println!("x: {}", x);
}
```

# Loops

- `for` is the most different from most C-like languages
  - `for` loops use an *iterator expression*:
  - `n..m` creates an iterator from `n` to `m` (exclusive).
  - Some data structures can be used as iterators, like arrays and `Vecs`.

```
// Loops from 0 to 9.
```

```
for x in 0..10 {
 println!("{}", x);
}
```

```
let xs = [0, 1, 2, 3, 4];
```

```
// Loop through elements in a slice of `xs`.
```

```
for x in &xs {
 println!("{}", x);
}
```



# Functions

```
fn foo(x: T, y: U, z: V) -> T {
 // ...
}
```

- `foo` is a function that takes three parameters:
  - `x` of type `T`
  - `y` of type `U`
  - `z` of type `V`
- `foo` returns a `T`.
- Must explicitly define argument and return types.
  - The compiler is actually smart enough to figure this out for you, but Rust's designers decided it was better practice to force explicit function typing.

# Functions

- The final expression in a function is its return value.
  - Use `return` for *early* returns from a function.

```
fn square(n: i32) -> i32 {
 n * n
}
```

```
fn squareish(n: i32) -> i32 {
 if n < 5 { return n; }
 n * n
}
```

```
fn square_bad(n: i32) -> i32 {
 n * n;
}
```

- The last one won't even compile!
  - Why? It ends in a semicolon, so it evaluates to `()`.

# Function Objects

- Several things can be used as function objects:
  - Function pointers (a reference to a normal function)
  - Closures (more advanced)
- Much more straightforward than C function pointers:

```
let x: fn(i32) -> i32 = square;
```

- Can be passed by reference:

```
fn apply_twice(f: &Fn(i32) -> i32, x: i32) -> i32 {
 f(f(x))
}
```

```
// ...
```

```
let y = apply_twice(&square, 5);
```

# Macros!

- Macros are like functions, but they're named with `!` at the end.
- Can do generally very powerful stuff.
  - They actually generate code at compile time!
- Call and use macros like functions.
- You can define your own with `macro_rules!` `macro_name` blocks.
  - These are *very* complicated.
- Because they're so powerful, a lot of common utilities are defined as macros.

# print! & println!

- Print stuff out. Yay.
- Use `{ }` for general string interpolation, and `{ :? }` for debug printing.
  - Some types can only be printed with `{ :? }`, like arrays and `Vecs`.

```
print!("{} , {} , {} ", "foo", 3, true);
```

```
// => foo, 3, true
```

```
println!("{:?}", "{:?}", "foo", [1, 2, 3]);
```

```
// => "foo", [1, 2, 3]
```

# format!

- Uses `println!`-style string interpolation to create formatted `Strings`.

```
let fmted = format!("{}", {:x}, {:?} ", 12, 155,
Some("Hello"));
// fmted == "12, 9b, Some("Hello") "
```

# panic!(msg)

- Exits current task with given message.
- Don't do this lightly! It is better to handle and report errors explicitly.

```
if x < 0 {
 panic!("Oh noes!");
}
```

# assert! & assert\_eq!

- `assert!(condition)` panics if condition is false.
- `assert_eq!(left, right)` panics if `left != right`.
- Useful for testing and catching illegal conditions.

```
#[test]
fn test_something() {
 let actual = 1 + 2;
 assert!(actual == 3);
 assert_eq!(3, actual);
}
```



# unreachable!()

- Used to indicate that some code should not be reached.
- `panic!`s when reached.
- Can be useful to track down unexpected bugs (e.g. optimization bugs).

```
if false {
 unreachable!();
}
```

# unimplemented!()

- Shorthand for `panic!("not yet implemented")`
- You'll probably see this in your assignments a lot!

```
fn sum(x: Vec<i32>) -> i32 {
 // TODO
 unimplemented!();
}
```

# Match statements

```
let x = 3;
```

```
match x {
 1 => println!("one fish"), // <- comma required
 2 => {
 println!("two fish");
 println!("two fish");
 }, // <- comma optional when using braces
 _ => println!("no fish for you"), // "otherwise" case
}
```

- `match` takes an expression (`x`) and branches on a list of `value => expression` statements.
- The entire match evaluates to one expression.
  - Like `if`, all arms must evaluate to the same type.
- `_` is commonly used as a catch-all (cf. Haskell, OCaml).

# Match statements

```
let x = 3;
let y = -3;

match (x, y) {
 (1, 1) => println!("one"),
 (2, j) => println!("two, {}", j),
 (_, 3) => println!("three"),
 (i, j) if i > 5 && j < 0 => println!("On guard!"),
 (_, _) => println!(":<"),
}
```

- The matched expression can be any expression (l-value), including tuples and function calls.
  - Matches can bind variables. `_` is a throw-away variable name.
- You *must* write an exhaustive match in order to compile.
- Use `if`-guards to constrain a match to certain conditions.
- Patterns can get very complex, as we'll see later.

# Tooling

# Cargo

- Rust's package manager & build tool
- Create a new project:
  - `cargo new project_name` (library)
  - `cargo new project_name --bin` (executable)
- Build your project: `cargo build`
- Run your tests: `cargo test`
- Magic, right? How does this work?

# Cargo.toml

- Cargo uses the `Cargo.toml` file to declare and manage dependencies and project metadata.
  - TOML is a simple format similar to INI.

```
[package]
name = "Rust"
version = "0.1.0"
authors = ["Ferris <ferris@rust-lang.org>"]
```

```
[dependencies]
uuid = "0.1"
rand = "0.3"
```

```
[profile.release]
opt-level = 3
debug = false
```

# cargo test

- A test is any function annotated with `#[test]`.
- `cargo test` will run all annotated functions in your project.
- Any function which executes without crashing (`panic!`ing) succeeds.
- Use `assert!` (or `assert_eq!`) to check conditions (and `panic!` on failure)

```
#[test]
fn it_works() {
 // ...
}
```



# Activity interlude

# Ownership & Borrowing

# Ownership & Borrowing

- Explicit ownership is the biggest new feature that Rust brings to the table!
- Ownership is all<sup>1</sup> checked at compile time!
- Newcomers to Rust often find themselves “fighting with the borrow checker” trying to get their code to compile

<sup>1</sup>*mostly*

# Ownership

- A variable binding *takes ownership* of its data.
  - A piece of data can only have one owner at a time.
- When a binding goes out of scope, the bound data is released automatically.
  - For heap-allocated data, this means de-allocation.
- Data *must be guaranteed* to outlive its references.

```
fn foo() {
 // Creates a Vec object.
 // Gives ownership of the Vec object to v1.
 let mut v1 = vec![1, 2, 3];

 v1.pop();
 v1.push(4);

 // At the end of the scope, v1 goes out of scope.
 // v1 still owns the Vec object, so it can be cleaned up.
}
```

# Move Semantics

```
let v1 = vec![1, 2, 3];
```

```
// Ownership of the Vec object moves to v2.
```

```
let v2 = v1;
```

```
println!("{}", v1[2]); // error: use of moved value `v1`
```

- `let v2 = v1;`
  - We don't want to copy the data, since that's expensive.
  - The data cannot have multiple owners.
  - Solution: move the Vec's ownership into `v2`, and declare `v1` invalid.
- `println!("{}", v1[2]);`
  - We know that `v1` is no longer a valid variable binding, so this is an error.
- Rust can reason about this at compile time, so it throws a compiler error.

# Move Semantics

- Moving ownership is a compile-time semantic; it doesn't involve moving data during your program.
- Moves are automatic (via assignments); no need to use something like C++'s `std::move`.
  - However, there are functions like `std::mem::replace` in Rust to provide advanced ownership management.

# Ownership

- Ownership does not always have to be moved.
- What would happen if it did? Rust would get very tedious to write:

```
fn vector_length(v: Vec<i32>) -> Vec<i32> {
 // Do whatever here,
 // then return ownership of `v` back to the caller
}
```

- You could imagine that this does not scale well either.
  - The more variables you had to hand back, the longer your return type would be!
  - Imagine having to pass ownership around for 5+ variables at a time :(

# Borrowing

- Obviously, this is not the case.
- Instead of transferring ownership, we can *borrow* data.
- A variable's data can be borrowed by taking a reference to the variable; ownership doesn't change.
  - When a reference goes out of scope, the borrow is over.
  - The original variable retains ownership throughout.

```
let v = vec![1, 2, 3];
```

```
// v_ref is a reference to v.
```

```
let v_ref = &v;
```

```
// use v_ref to access the data in the vector v.
```

```
assert_eq!(v[1], v_ref[1]);
```



# Borrowing

- Caveat: this adds restrictions to the original variable.
- Ownership cannot be transferred from a variable while references to it exist.
  - That would invalidate the reference.

```
let v = vec![1, 2, 3];
```

```
// v_ref is a reference to v.
```

```
let v_ref = &v;
```

```
// Moving ownership to v_new would invalidate v_ref.
```

```
// error: cannot move out of `v` because it is borrowed
```

```
let v_new = v;
```

# Borrowing

```
/// `length` only needs `vector` temporarily, so it is borrowed.
fn length(vec_ref: &Vec<i32>) -> usize {
 // vec_ref is auto-dereferenced when you call methods on it.
 vec_ref.len()
 // you can also explicitly dereference.
 // (*vec_ref).len()
}

fn main() {
 let vector = vec![];
 length(&vector);
 println!("{:?}", vector); // this is fine
}
```

- Note the type of `length`: `vec_ref` is passed by reference, so it's now an `&Vec<i32>`.
- References, like bindings, are *immutable* by default.
- The borrow is over after the reference goes out of scope (at the end of `length`).

# Borrowing

```
/// `push` needs to modify `vector` so it is borrowed mutably.
fn push(vec_ref: &mut Vec<i32>, x: i32) {
 vec_ref.push(x);
}

fn main() {
 let mut vector: Vec<i32> = vec![];
 let vector_ref: &mut Vec<i32> = &mut vector;
 push(vector_ref, 4);
}
```

- Variables can be borrowed by *mutable* reference: `&mut vec_ref`.
  - `vec_ref` is a reference to a mutable `Vec`.
  - The type is `&mut Vec<i32>`, not `&Vec<i32>`.
- Different from a reference which is variable.

# Borrowing

```
/// `push` needs to modify `vector` so it is borrowed mutably.
fn push2(vec_ref: &mut Vec<i32>, x: i32) {
 // error: cannot move out of borrowed content.
 let vector = *vec_ref;
 vector.push(x);
}

fn main() {
 let mut vector = vec![];
 push2(&mut vector, 4);
}
```

- Error! You can't dereference `vec_ref` into a variable binding because that would change the ownership of the data.

# Borrowing

- Rust will auto-dereference variables...
  - When making method calls on a reference.
  - When passing a reference as a function argument.

```
/// `length` only needs `vector` temporarily, so it is borrowed.
fn length(vec_ref: &&Vec<i32>) -> usize {
 // vec_ref is auto-dereferenced when you call methods on it.
 vec_ref.len()
}
```

```
fn main() {
 let vector = vec![];
 length(&&&&&&&&&vector);
}
```

# Borrowing

- You will have to dereference variables...
  - When writing into them.
  - And other times that usage may be ambiguous.

```
let mut a = 5;
let ref_a = &mut a;
*ref_a = 4;
println!("{}", *ref_a + 4);
// ==> 8
```

# Copy Types

- Rust defines a trait<sup>1</sup> named `Copy` that signifies that a type may be copied instead whenever it would be moved.
- Most primitive types are `Copy` (`i32`, `f64`, `char`, `bool`, etc.)
- Types that contain references may not be `Copy` (e.g. `Vec`, `String`).

```
let x: i32 = 12;
let y = x; // `i32` is `Copy`, so it's not moved :D
println!("x still works: {}, and so does y: {}", x, y);
```

<sup>1</sup> Like a Java interface or Haskell typeclass

# Borrowing Rules

## *The Holy Grail of Rust*

Learn these rules, and they will serve you well.

You can't keep borrowing something after it stops existing.

One object may have many immutable references to it (`&T`).

**OR** *exactly one* mutable reference (`&mut T`) (not both).

That's it!



# Borrowing Prevents...

- Iterator invalidation due to mutating a collection you're iterating over.
- This pattern can be written in C, C++, Java, Python, Javascript...
  - But may result in, e.g, `ConcurrentModificationException` (at runtime!)

```
let mut vs = vec![1,2,3,4];
for v in &vs {
 vs.pop();
 // ERROR: cannot borrow `vs` as mutable because
 // it is also borrowed as immutable
}
```

- `pop` needs to borrow `vs` as mutable in order to modify the data.
- But `vs` is being borrowed as immutable by the loop!

# Borrowing Prevents...

- Use-after-free
- Valid in C, C++...

```
let y: &i32;
{
 let x = 5;
 y = &x; // error: `x` does not live long enough
}
println!("{}", *y);
```

- The full error message:

```
error: `x` does not live long enough
note: reference must be valid for the block suffix following statement
 0 at 1:16
...but borrowed value is only valid for the block suffix
 following statement 0 at 4:18
```

- This eliminates a *huge* number of memory safety bugs *at compile time*.

# Example: Vectors

- You can iterate over Vecs in three different ways:

```
let mut vs = vec![0,1,2,3,4,5,6];
```

```
// Borrow immutably
```

```
for v in &vs { // Can also write `for v in vs.iter()`
 println!("I'm borrowing {}. ", v);
}
```

```
// Borrow mutably
```

```
for v in &mut vs { // Can also write `for v in vs.iter_mut()`
 *v = *v + 1;
 println!("I'm mutably borrowing {}. ", v);
}
```

```
// Take ownership of the whole vector
```

```
for v in vs { // Can also write `for v in vs.into_iter()`
 println!("I now own {}! AHAHAHAHA!", v);
}
```

```
// `vs` is no longer valid
```

# Looking ahead

- Assignment 1 now due 10p Feb 14
- Assignment 2 will be out right after
  - Read it before coming to class
  - It's a longer assignment, but you have more time to do it
  - Use the discussion board for this assignment
- Make sure you submit Lab 1 to Blackboard
- Read the Rust book chapters for next time
  - Follow along by typing code into your Rust environment

## Lesson objectives

- Become familiar with basic Rust syntax and types
- Apply Rust languages features when writing code
- Intuitively understand Rust's borrowing and ownership semantics